

Excercise 1 - Data Exploration and Vizualization

February 17, 2025

1 Data Quality Check

Fetch - Data Analyst Take Home Created by: Sakshi Rajendra Khadayate

```
[1]: # Import necessary libraries
import pandas as pd
import unicodedata
```

```
[2]: # Import CSV files using pandas
# Encoding is set to 'utf-8' to handle any special characters that may be
    ↪ present in the files
df_users = pd.read_csv('USER TAKEHOME.csv', encoding='utf-8')
df_products = pd.read_csv('PRODUCTS TAKEHOME.csv', encoding = 'utf-8')
df_transactions = pd.read_csv('TRANSACTION TAKEHOME (1).csv', encoding='utf-8')
```

```
[3]: df_users.head()
```

```
[3]:
```

	ID	CREATED_DATE	\
0	5ef3b4f17053ab141787697d	2020-06-24 20:17:54.000	Z
1	5ff220d383fcfc12622b96bc	2021-01-03 19:53:55.000	Z
2	6477950aa55bb77a0e27ee10	2023-05-31 18:42:18.000	Z
3	658a306e99b40f103b63ccf8	2023-12-26 01:46:22.000	Z
4	653cf5d6a225ea102b7ecdc2	2023-10-28 11:51:50.000	Z

	BIRTH_DATE	STATE	LANGUAGE	GENDER
0	2000-08-11 00:00:00.000	Z CA	es-419	female
1	2001-09-24 04:00:00.000	Z PA	en	female
2	1994-10-28 00:00:00.000	Z FL	es-419	female
3	NaN	NC	en	NaN
4	1972-03-19 00:00:00.000	Z PA	en	female

```
[4]: df_products.head()
```

```
[4]:
```

	CATEGORY_1	CATEGORY_2	CATEGORY_3	\
0	Health & Wellness	Sexual Health	Conductivity Gels & Lotions	
1	Snacks	Puffed Snacks	Cheese Curls & Puffs	
2	Health & Wellness	Hair Care	Hair Care Accessories	
3	Health & Wellness	Oral Care	Toothpaste	

4 Health & Wellness Medicines & Treatments Essential Oils

	CATEGORY_4	MANUFACTURER \
0	NaN	NaN
1	NaN	NaN
2	NaN	PLACEHOLDER MANUFACTURER
3	NaN	COLGATE-PALMOLIVE
4	NaN	MAPLE HOLISTICS AND HONEYDEW PRODUCTS INTERCHA...

	BRAND	BARCODE
0	NaN	7.964940e+11
1	NaN	2.327801e+10
2	ELECSOP	4.618180e+11
3	COLGATE	3.500047e+10
4	MAPLE HOLISTICS	8.068110e+11

```
[5]: df_transactions.head()
```

	RECEIPT_ID	PURCHASE_DATE \
0	0000d256-4041-4a3e-adc4-5623fb6e0c99	2024-08-21
1	0001455d-7a92-4a7b-a1d2-c747af1c8fd3	2024-07-20
2	00017e0a-7851-42fb-bfab-0baa96e23586	2024-08-18
3	000239aa-3478-453d-801e-66a82e39c8af	2024-06-18
4	00026b4c-dfe8-49dd-b026-4c2f0fd5c6a1	2024-07-04

	SCAN_DATE	STORE_NAME	USER_ID \
0	2024-08-21 14:19:06.539 Z	WALMART	63b73a7f3d310dceeabd4758
1	2024-07-20 09:50:24.206 Z	ALDI	62c08877baa38d1a1f6c211a
2	2024-08-19 15:38:56.813 Z	WALMART	60842f207ac8b7729e472020
3	2024-06-19 11:03:37.468 Z	FOOD LION	63fcd7cea4f8442c3386b589
4	2024-07-05 15:56:43.549 Z	RANDALLS	6193231ae9b3d75037b0f928

	BARCODE	FINAL_QUANTITY	FINAL_SALE
0	1.530001e+10	1.00	
1	NaN	zero	1.49
2	7.874223e+10	1.00	
3	7.833997e+11	zero	3.49
4	4.790050e+10	1.00	

```
[6]: # Function to check and fix duplicates and null values for a given primary key
def remove_null_and_duplicate(df, primary_key):

    # Print the number of rows and columns
    print("Shape of DataFrame:", df.shape)

    # Print count of missing values in each column
    print("\nMissing Values:\n", df.isnull().sum())
```

```

# Print the number of unique values in each column
print("\nUnique Values Count:\n", df.nunique())

# Drop rows where primary key column has missing (NaN) values
df = df.dropna(subset=[primary_key])

# Check if primary key column has duplicate values
if df[primary_key].duplicated().any():
    print(f"\nWarning: Duplicate values found in '{primary_key}'. Keeping
↳only the first occurrence.")

# Keep only the first occurrence of each duplicate
df = df.drop_duplicates(subset=[primary_key])

return df

```

```

[7]: # Run the function to remove_null_and_duplicate (Users and Product tables has
↳primary keys)
print ("Users")
df_users = remove_null_and_duplicate(df_users, 'ID')
print("-----")
print ("Products")
df_products = remove_null_and_duplicate(df_products, 'BARCODE')
print("-----")

print ("Transactions")
# Print the number of rows and columns
print("Shape of DataFrame:", df_transactions.shape)

# Print count of missing values in each column
print("\nMissing Values:\n", df_transactions.isnull().sum())

# Print the number of unique values in each column
print("\nUnique Values Count:\n", df_transactions.nunique())

```

Users

Shape of DataFrame: (100000, 6)

Missing Values:

ID	0
CREATED_DATE	0
BIRTH_DATE	3675
STATE	4812
LANGUAGE	30508
GENDER	5892

dtype: int64

Unique Values Count:

ID	100000
CREATED_DATE	99942
BIRTH_DATE	54721
STATE	52
LANGUAGE	2
GENDER	11

dtype: int64

Products

Shape of DataFrame: (20773, 7)

Missing Values:

CATEGORY_1	2
CATEGORY_2	34
CATEGORY_3	1554
CATEGORY_4	19186
MANUFACTURER	5547
BRAND	5547
BARCODE	112

dtype: int64

Unique Values Count:

CATEGORY_1	14
CATEGORY_2	53
CATEGORY_3	202
CATEGORY_4	51
MANUFACTURER	1723
BRAND	3191
BARCODE	13590

dtype: int64

Warning: Duplicate values found in 'BARCODE'. Keeping only the first occurrence.

Transactions

Shape of DataFrame: (50000, 8)

Missing Values:

RECEIPT_ID	0
PURCHASE_DATE	0
SCAN_DATE	0
STORE_NAME	0
USER_ID	0
BARCODE	5762
FINAL_QUANTITY	0
FINAL_SALE	0

dtype: int64

```

Unique Values Count:
RECEIPT_ID      24440
PURCHASE_DATE   89
SCAN_DATE       24440
STORE_NAME      954
USER_ID         17694
BARCODE         11027
FINAL_QUANTITY  87
FINAL_SALE      1435
dtype: int64

```

```

[8]: # Trim spaces from column names to ensure no extra spaces cause issues during
      ↪data manipulation
df_users.columns = df_users.columns.str.strip()
df_products.columns = df_products.columns.str.strip()
df_transactions.columns = df_transactions.columns.str.strip()

# Trim spaces from all string values in the DataFrame to clean any leading or
      ↪trailing spaces in the actual data
df_users = df_users.applymap(lambda x: x.strip() if isinstance(x, str) else x)
df_products = df_products.applymap(lambda x: x.strip() if isinstance(x, str)
      ↪else x)
df_transactions = df_transactions.applymap(lambda x: x.strip() if isinstance(x,
      ↪str) else x)

```

```

[9]: # Function to check if a string contains non-ASCII characters
def contains_non_ascii(s):
    if isinstance(s, str): # Check if it's a string
        return not all(ord(c) < 128 for c in s)
    return False # Non-strings are considered valid as they don't contain
      ↪non-ASCII characters

# Function to check all string columns for non-ASCII values
def find_non_ascii_rows(df):

    # Select only string columns
    string_cols = df.select_dtypes(include=['object'])

    # Identify non-ASCII values
    mask = string_cols.applymap(contains_non_ascii)

    # Get rows where any string column has a non-ASCII value
    non_ascii_rows = df[mask.any(axis=1)]

    # Print the column names where non-ASCII values are found
    if not non_ascii_rows.empty:
        affected_columns = mask.any(axis=0)

```

```

        print("Non-ASCII values found in columns in :",
↪list(affected_columns[affected_columns].index))
    else:
        print("No Non-ASCII values found")

    return non_ascii_rows

```

```

[10]: # Find and print rows with non-ASCII characters
print ("Users")
non_ascii_rows = find_non_ascii_rows(df_users)
print ("\nProducts")
non_ascii_rows = find_non_ascii_rows(df_products)
print ("\nTransactions")
non_ascii_rows = find_non_ascii_rows(df_transactions)

```

Users

No Non-ASCII values found

Products

Non-ASCII values found in columns in : ['MANUFACTURER', 'BRAND']

Transactions

Non-ASCII values found in columns in : ['STORE_NAME']

```

[11]: # =====
# Users Table Data Cleaning and Transformation
# =====

# Convert Date column to SQL-friendly format ('YYYY-MM-DD HH:MM:SS')
df_users['CREATED_DATE'] = pd.to_datetime(df_users['CREATED_DATE']).dt.
↪strftime('%Y-%m-%d %H:%M:%S')
df_users['BIRTH_DATE'] = pd.to_datetime(df_users['BIRTH_DATE']).dt.
↪strftime('%Y-%m-%d %H:%M:%S')

# Fill missing 'BIRTH_DATE' values with a default pseudo date ('1900-01-01 00:
↪00:00'), making the data consistent
df_users['BIRTH_DATE'] = df_users['BIRTH_DATE'].fillna('1900-01-01 00:00:00')

# Save the updated file
output_file = "converted_file_users.csv"
df_users.to_csv(output_file, index=False)

print(f"Successfully converted file in desired format, Saved as: {output_file}")

```

Successfully converted file in desired format, Saved as:

converted_file_users.csv

```
[12]: # =====
# Products Table Data Cleaning and Transformation
# =====

# Function to replace non-ASCII characters in strings
def replace_non_ascii(s):
    if isinstance(s, str): # Apply only to strings
        return unicodedata.normalize('NFKD', s).encode('ascii', 'ignore').
        ↪decode('ascii')
    return s # Return non-string values as they are

# Select only string columns
string_cols = df_products.select_dtypes(include=['object'])

# Apply the replacement function to only string columns
df_products[string_cols.columns] = string_cols.applymap(replace_non_ascii)

# Save the updated file
output_file = "converted_file_products.csv"
df_products.to_csv(output_file, index=False, encoding='utf-8')

print(f"Successfully converted file in desired format, Saved as: {output_file}")
```

Successfully converted file in desired format, Saved as:
converted_file_products.csv

```
[13]: # =====
# Transactions Table Data Cleaning and Transformation
# =====

# # Convert Date column to SQL-friendly format ('YYYY-MM-DD HH:MM:SS'), ↵
↪replace blanks with 0
df_transactions['PURCHASE_DATE'] = pd.
    ↪to_datetime(df_transactions['PURCHASE_DATE']).dt.strftime('%Y-%m-%d %H:%M:
    ↪%S')
df_transactions['SCAN_DATE'] = pd.to_datetime(df_transactions['SCAN_DATE']).dt.
    ↪strftime('%Y-%m-%d %H:%M:%S')
df_transactions['BARCODE'] = df_transactions['BARCODE'].fillna(0)
df_transactions = df_transactions[df_transactions['FINAL_QUANTITY']!= 'zero']
df_transactions = df_transactions[df_transactions['FINAL_SALE']!= '']

# Function to replace non-ASCII characters in strings
def replace_non_ascii(s):
    if isinstance(s, str): # Apply only to strings
        return unicodedata.normalize('NFKD', s).encode('ascii', 'ignore').
        ↪decode('ascii')
    return s # Return non-string values as they are
```

```

# Select only string columns
string_cols = df_transactions.select_dtypes(include=['object'])

# Apply the replacement function to only string columns
df_transactions[string_cols.columns] = string_cols.applymap(replace_non_ascii)

# Save the updated file
output_file = "converted_file_transactions.csv"
df_transactions.to_csv(output_file, index=False)

print(f"Successfully converted file in desired format, Saved as: {output_file}")

```

Successfully converted file in desired format, Saved as:
converted_file_transactions.csv

```
[14]: df_users.head()
```

```

[14]:
          ID          CREATED_DATE          BIRTH_DATE STATE \
0  5ef3b4f17053ab141787697d  2020-06-24 20:17:54  2000-08-11 00:00:00  CA
1  5ff220d383fcfc12622b96bc  2021-01-03 19:53:55  2001-09-24 04:00:00  PA
2  6477950aa55bb77a0e27ee10  2023-05-31 18:42:18  1994-10-28 00:00:00  FL
3  658a306e99b40f103b63ccf8  2023-12-26 01:46:22  1900-01-01 00:00:00  NC
4  653cf5d6a225ea102b7ecdc2  2023-10-28 11:51:50  1972-03-19 00:00:00  PA

LANGUAGE  GENDER
0    es-419  female
1         en  female
2    es-419  female
3         en     NaN
4         en  female

```

```
[15]: df_products.head()
```

```

[15]:
          CATEGORY_1          CATEGORY_2          CATEGORY_3 \
0  Health & Wellness  Sexual Health  Conductivity Gels & Lotions
1          Snacks      Puffed Snacks      Cheese Curls & Puffs
2  Health & Wellness      Hair Care      Hair Care Accessories
3  Health & Wellness      Oral Care      Toothpaste
4  Health & Wellness  Medicines & Treatments      Essential Oils

CATEGORY_4          MANUFACTURER \
0         NaN         NaN
1         NaN         NaN
2         NaN  PLACEHOLDER MANUFACTURER
3         NaN      COLGATE-PALMOLIVE
4         NaN  MAPLE HOLISTICS AND HONEYDEW PRODUCTS INTERCHA...

```


	BRAND	BARCODE
0	NaN	7.964940e+11
1	NaN	2.327801e+10
2	ELECSOP	4.618180e+11
3	COLGATE	3.500047e+10
4	MAPLE HOLISTICS	8.068110e+11

```
[16]: df_transactions.head()
```

```
[16]:
```

	RECEIPT_ID	PURCHASE_DATE	\
25000	7b3ec72d-9d30-40b8-b185-0bfb638942a9	2024-08-20 00:00:00	
25001	04869b68-29e3-4e8d-9bdb-950046fc3473	2024-08-05 00:00:00	
25002	f1a96308-24a5-46a8-8d8c-285cf9dce1ba	2024-09-03 00:00:00	
25003	7ee1798e-fd2e-4278-838b-f417fdcafe08	2024-08-30 00:00:00	
25004	21feab39-49f2-42e9-ae69-10371e2fc0a9	2024-08-23 00:00:00	

	SCAN_DATE	STORE_NAME	USER_ID	\
25000	2024-08-20 11:17:29	DOLLAR GENERAL STORE	60fc1e6deb7585430ff52ee7	
25001	2024-08-09 16:06:00	DOLLAR GENERAL STORE	654cf234a225ea102b81072e	
25002	2024-09-03 11:28:25	WALMART	63c1cb6d3d310dceeac55487	
25003	2024-09-04 12:53:31	DOLLAR GENERAL STORE	65c29b137050d0a6206cd24f	
25004	2024-08-27 10:45:00	TARGET	61a58ac49c135b462ccddd1c	

	BARCODE	FINAL_QUANTITY	FINAL_SALE
25000	7.455271e+11	1.00	1.65
25001	7.455271e+11	1.00	1.65
25002	3.700083e+10	1.00	28.22
25003	1.200050e+10	1.00	5.25
25004	2.400039e+10	1.00	2.59

----- All 3 files are ready to be imported in SQL for further Analysis -----

2

3 Data Visualization

```
[17]: import pandas as pd
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots
import numpy as np
from datetime import datetime

# Read the CSV files
df_users = pd.read_csv('converted_file_users.csv')
```

```

df_products = pd.read_csv('converted_file_products.csv')
df_transactions = pd.read_csv('converted_file_transactions.csv')

#Convert date columns to datetime
df_users['CREATED_DATE'] = pd.to_datetime(df_users['CREATED_DATE'])
df_users['BIRTH_DATE'] = pd.to_datetime(df_users['BIRTH_DATE'])
df_transactions['PURCHASE_DATE'] = pd.
    ↳to_datetime(df_transactions['PURCHASE_DATE'])
df_transactions['SCAN_DATE'] = pd.to_datetime(df_transactions['SCAN_DATE'])

# 1. User Language Distribution
def plot_language_distribution():
    lang_dist = df_users['LANGUAGE'].value_counts().reset_index()
    lang_dist.columns = ['Language', 'Count']

    fig = px.pie(lang_dist,
                  values='Count',
                  names='Language',
                  title='User Language Distribution',
                  color_discrete_sequence=px.colors.qualitative.Set3)

    fig.update_traces(textposition='inside', textinfo='percent+label')
    fig.show()

# 2. User Gender Distribution
def plot_gender_distribution():
    lang_dist = df_users['GENDER'].value_counts().reset_index()
    lang_dist.columns = ['Gender', 'Count']

    fig = px.pie(lang_dist,
                  values='Count',
                  names='Gender',
                  title='User Gender Distribution', # Corrected title
                  color_discrete_sequence=px.colors.qualitative.Set3)

    fig.update_traces(textposition='inside', textinfo='percent+label')
    fig.show()

# 3. User Age Distribution
def plot_age_distribution():
    df_users['AGE'] = (datetime.now() - df_users['BIRTH_DATE']).dt.days / 365.25

    fig = px.histogram(df_users[df_users['AGE'] < 100], # Filter out
    ↳unrealistic ages
                      x='AGE',
                      nbins=50,
                      title='User Age Distribution',

```

```

        color_discrete_sequence=['#FF7F50'])

fig.update_layout(xaxis_title='Age',
                  yaxis_title='Count')
fig.show()

# 4. Top Stores Analysis
def plot_top_stores():
    store_counts = df_transactions['STORE_NAME'].value_counts().head(10)

    fig = px.bar(x=store_counts.index,
                 y=store_counts.values,
                 title='Top 10 Stores by Transaction Count',
                 color_discrete_sequence=['#4169E1'])

    fig.update_layout(xaxis_title='Store Name',
                      yaxis_title='Number of Transactions',
                      xaxis_tickangle=45)

    fig.show()

# 5. Category Analysis
def plot_category_analysis():
    # Create subplots for different category levels (3 rows, 2 columns)
    fig = make_subplots(rows=3, cols=2,
                        subplot_titles=('Category 1', 'Category 2', 'Category 3',
                                       'Category 4', 'Top Brands', 'Top_
↳Manufacturers'))

    # Category 1 Distribution
    cat1_counts = df_products['CATEGORY_1'].value_counts().head(10)
    fig.add_trace(go.Bar(x=cat1_counts.index, y=cat1_counts.values),
                  row=1, col=1)

    # Category 2 Distribution
    cat2_counts = df_products['CATEGORY_2'].value_counts().head(10)
    fig.add_trace(go.Bar(x=cat2_counts.index, y=cat2_counts.values),
                  row=1, col=2)

    # Category 3 Distribution
    cat3_counts = df_products['CATEGORY_3'].value_counts().head(10)
    fig.add_trace(go.Bar(x=cat3_counts.index, y=cat3_counts.values),
                  row=2, col=1)

    # Category 4 Distribution
    cat4_counts = df_products['CATEGORY_4'].value_counts().head(10)
    fig.add_trace(go.Bar(x=cat4_counts.index, y=cat4_counts.values),
                  row=2, col=2)

```

```

# Top Brands (Top 10)
brand_counts = df_products['BRAND'].value_counts().head(10)
fig.add_trace(go.Bar(x=brand_counts.index, y=brand_counts.values),
               row=3, col=1)

# Top Manufacturers (Top 10)
manufacturer_counts = df_products['MANUFACTURER'].value_counts().head(10)
fig.add_trace(go.Bar(x=manufacturer_counts.index, y=manufacturer_counts.
↪values),
               row=3, col=2)

# Update the layout
fig.update_layout(height=1000,
                  showlegend=False,
                  title_text="Product Category Analysis")

# Show the plot
fig.show()

# 6. Transaction Trends Over Time
def plot_transaction_trends():
    daily_transactions = df_transactions.groupby('PURCHASE_DATE').agg({
        'RECEIPT_ID': 'count',
        'FINAL_SALE': 'sum'
    }).reset_index()

    # Create figure with secondary y-axis
    fig = make_subplots(specs=[[{"secondary_y": True}]])

    fig.add_trace(
        go.Scatter(x=daily_transactions['PURCHASE_DATE'],
                   y=daily_transactions['RECEIPT_ID'],
                   name="Number of Transactions"),
        secondary_y=False,
    )

    fig.add_trace(
        go.Scatter(x=daily_transactions['PURCHASE_DATE'],
                   y=daily_transactions['FINAL_SALE'],
                   name="Total Sales"),
        secondary_y=True,
    )

    fig.update_layout(

```

```

        title_text="Daily Transaction Trends",
        xaxis_title="Date"
    )

    fig.update_yaxes(title_text="Number of Transactions", secondary_y=False)
    fig.update_yaxes(title_text="Total Sales ($)", secondary_y=True)

    fig.show()

# 7. Days Between Purchase and Scan Dates
def plot_purchase_scan_diff():
    # Calculate the difference in days between PURCHASE_DATE and SCAN_DATE
    df_transactions['DIFFERENCE_DAYS'] = (df_transactions['SCAN_DATE'] -
    ↪df_transactions['PURCHASE_DATE']).dt.days

    # Filter out negative or unrealistic differences (optional, depending on
    ↪the dataset)
    df_transactions_filtered =
    ↪df_transactions[df_transactions['DIFFERENCE_DAYS'] >= 0]

    # Calculate the average difference in days
    avg_days = df_transactions_filtered['DIFFERENCE_DAYS'].mean()

    # Plot the histogram of the difference in days
    fig = px.histogram(df_transactions_filtered,
                        x='DIFFERENCE_DAYS',
                        nbins=50,
                        title='Distribution of Days Between Purchase and Scan
    ↪Dates',
                        color_discrete_sequence=['#1f77b4'])

    # Get the min and max values of the 'DIFFERENCE_DAYS' for better tick
    ↪control
    min_day = int(df_transactions_filtered['DIFFERENCE_DAYS'].min())
    max_day = int(df_transactions_filtered['DIFFERENCE_DAYS'].max())

    # Update the layout to show all x-axis labels and add the average line
    fig.update_layout(
        xaxis_title='Days Difference',
        yaxis_title='Count of Transactions',
        bargap=0.2,
        xaxis=dict(
            tickmode='array', # Set the x-axis ticks as an array
            tickvals=list(range(min_day, max_day + 1)), # Show ticks for each
    ↪day from min to max

```

```

        ticktext=list(map(str, range(min_day, max_day + 1))) # Display
↳ each tick value as a string
    )
)

# Add vertical line for the average
fig.add_vline(x=avg_days,
              line=dict(color='red', dash='dash'),
              annotation_text=f"Average: {avg_days:.2f} days",
              annotation_position="top right")

fig.show()

# 8. User State Distribution Map
def plot_state_distribution():
    state_counts = df_users['STATE'].value_counts().reset_index()
    state_counts.columns = ['state', 'count']

    fig = px.choropleth(state_counts,
                        locations='state',
                        locationmode="USA-states",
                        color='count',
                        scope="usa",
                        title="User Distribution by State",
                        color_continuous_scale="Viridis")

    fig.show()

# Execute all visualizations
def main():
    plot_language_distribution()
    plot_gender_distribution()
    plot_age_distribution()
    plot_top_stores()
    plot_category_analysis()
    plot_transaction_trends()
    plot_purchase_scan_diff()
    plot_state_distribution()
    print("All visualizations generated!")

if __name__ == "__main__":
    main()

```


User Language Distribution




User Gender Distribution



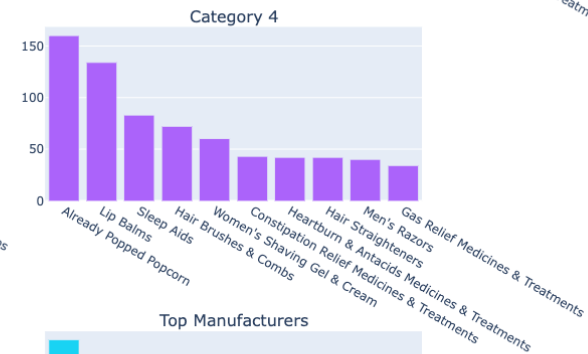

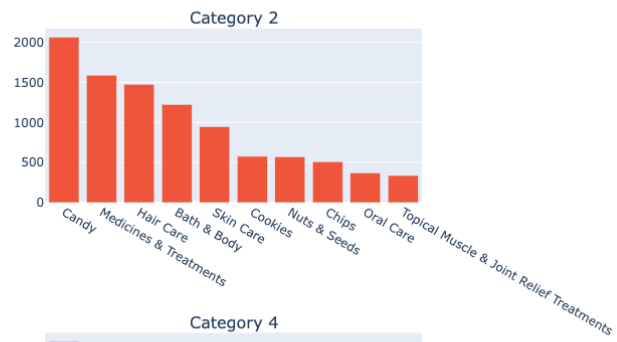



User Age Distribution







Product Category Analysis



Daily Transaction Trends



Distribution of Days Between Purchase and Scan Dates



User Distribution by State



All visualizations generated!

4